

Outline of this lecture

Collective Communications

- Types of collective communications
- Global synchronization
- Global Reduction

Collective Communications

General concepts

- Collective communications allow making a series of point-to-point communications in one single call.
- A collective communication always concerns all the processes of the indicated communicator.
- For each process, the call ends when its participation in the collective call is completed, in the sense of point-to-point communications (therefore, when the concerned memory area can be changed).
- The management of tags in these communications is transparent and system-dependent. Therefore, they are never explicitly defined during calls to subroutines. An advantage of this is that collective communications never interfere with point-to-point communications.

Collective Communications

Types of collective communications

- One which ensures global synchronizations : `MPI_BARRIER()`
- Ones which only transfer data :
 - Global distribution of data : `MPI_BCAST()`
 - Selective distribution of data : `MPI_SCATTER()`
 - Collection of distributed data : `MPI_GATHER()`
 - Collection of distributed data by all the processes : `MPI_ALLGATHER()`
 - Collection and selective distribution by all the processes of distributed data : `MPI_ALLTOALL()`
- Ones which, in addition to the communications management, carry out operations on the transferred data :
 - Reduction operations (sum, product, maximum, minimum, etc.), whether of a predefined or personal type : `MPI_REDUCE()`
 - Reduction operations with distributing of the result (this is in fact equivalent to an `MPI_REDUCE()` followed by an `MPI_BCAST()`) : `MPI_ALLREDUCE()`

Collective Communications

Global synchronization : MPI_BARRIER()

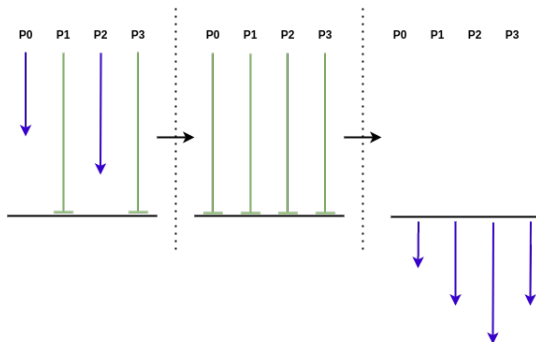


Figure: Global Synchronization : MPI_BARRIER()

1

```
COMM.Barrier()
```

Collective Communications

Global synchronization : `MPI_BCAST()`

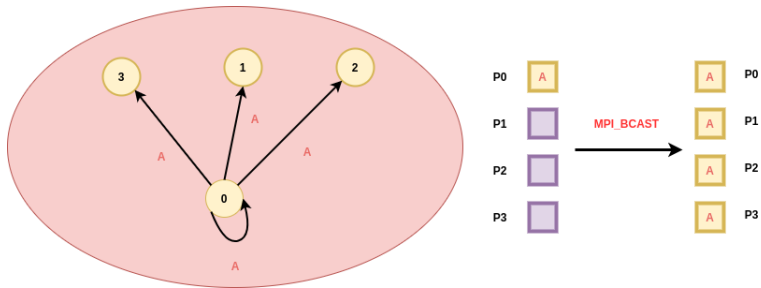


Figure: Global distribution : `MPI_BCAST()`

Collective Communications

Global synchronization : MPI_BCAST()

```
1 MPI_BCAST(buffer, count, datatype, root, comm, code)
2
3 <type> :: buffer
4 integer :: count, datatype, root, comm, code
```

- Send, starting at position buffer, a message of count element of type datatype, by the root process, to all the members of communicator comm.
- Receive this message at position buffer for all the processes other than the root.

```
1 Comm.bcast(self, obj, int root=0)
2 # or
3 Comm.Bcast(self, buf, int root=0)
```

Collective Communications

Global synchronization : MPI_BCAST() Full example 1

```
1  from mpi4py import MPI
2  import numpy as np
3
4  COMM = MPI.COMM_WORLD
5  RANK = COMM.Get_rank()
6
7  if RANK == 0:
8      sendbuf = {'key1' : [7, 2.72, 2+3j]}
9  else:
10     sendbuf = None
11
12  recvbuf = COMM.bcast(sendbuf, root=0)
13
14  print("I am the process {rank}, I received data {data} from 2".format(rank=RANK↔
    , data=recvbuf))
```

`mpirun -n 4 python bcast.py`

```
I am the process 0, I received data {'key1': [7, 2.72, (2+3j)]} from 2
I am the process 2, I received data {'key1': [7, 2.72, (2+3j)]} from 2
I am the process 1, I received data {'key1': [7, 2.72, (2+3j)]} from 2
I am the process 3, I received data {'key1': [7, 2.72, (2+3j)]} from 2
```

Collective Communications

Global synchronization : MPI_BCAST() Full example 2

```
1  from mpi4py import MPI
2  import numpy as np
3
4  COMM = MPI.COMM_WORLD
5  RANK = COMM.Get_rank()
6
7  if RANK == 0:
8      sendbuf = np.arange(10, dtype='i')
9  else:
10     sendbuf = np.empty(10, dtype='i')
11
12     COMM.Bcast(sendbuf, root=0)
13
14     print("I am the process {rank}, I received data {data} from 2".format(rank=RANK, data=sendbuf))
```

`mpirun -n 4 python bcast.py`

```
I am the process 1, I received data [0 1 2 3 4 5 6 7 8 9] from 2
I am the process 3, I received data [0 1 2 3 4 5 6 7 8 9] from 2
I am the process 0, I received data [0 1 2 3 4 5 6 7 8 9] from 2
I am the process 2, I received data [0 1 2 3 4 5 6 7 8 9] from 2
```


Collective Communications

Global synchronization : MPI_SCATTER()

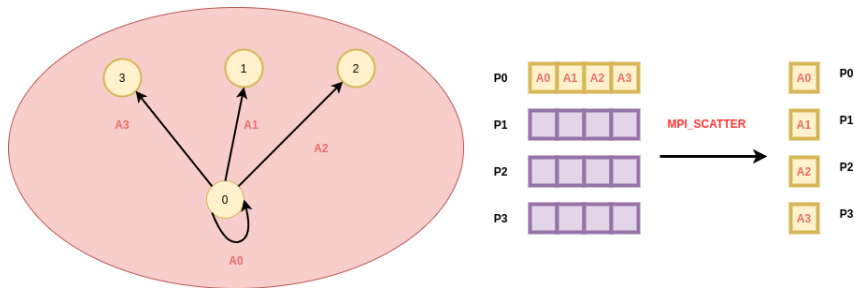


Figure: Selected distribution : MPI_SCATTER()

Collective Communications

Global synchronization : MPI_SCATTER()

```
1 MPI_SCATTER(message_a_repartir, longueur_message_emis, type_message_emis,
2             message_recu, longueur_message_recu, type_message_recu, rang_source, comm ←
3             , code)
4 <type> :: sendbuf, recvbuf
5 integer :: sendcount, recvcnt
6 integer :: sendtype, recvtype
7 integer :: root, comm, code
```

- Scatter by process root, starting at position sendbuf, message sendcount element of type sendtype, to all the processes of communicator comm.
- Receive this message at position recvbuf, of recvcnt element of type recvtype for all processes of communicator comm.
- ➡ The couples (sendcount, sendtype) and (recvcnt, recvtype) must represent the same quantity of data.
- ➡ Data are scattered in chunks of same size ; a chunk consists of sendcount elements of type sendtype.
- ➡ The i-th chunk is sent to the i-th process.

```
1 Comm.scatter(self, sendobj, int root=0)
2 #or
3 Comm.Scatter(self, sendbuf, recvbuf, int root=0)
```

Collective Communications

Global synchronization : MPI_SCATTER() Full example

```
1
2  from mpi4py import MPI
3  import numpy as np
4
5
6  COMM = MPI.COMM_WORLD
7  SIZE = COMM.Get_size()
8  RANK = COMM.Get_rank()
9
10 if RANK == 0:
11     sendbuf = [(i+1)**2 for i in range(SIZE)]
12     print("The data will be scattered is", sendbuf)
13 else:
14     sendbuf = None
15
16 recvbuf = COMM.scatter(sendbuf, root=0)
17 assert recvbuf == (RANK+1)**2
```

`mpirun -n 4 python scatter.py`

```
The data will be scattered is [1, 4, 9, 16]
I am the process 0, I received data 1 from 0
I am the process 3, I received data 16 from 0
I am the process 1, I received data 4 from 0
I am the process 2, I received data 9 from 0
```

Collective Communications

Global synchronization : MPI_GATHER()

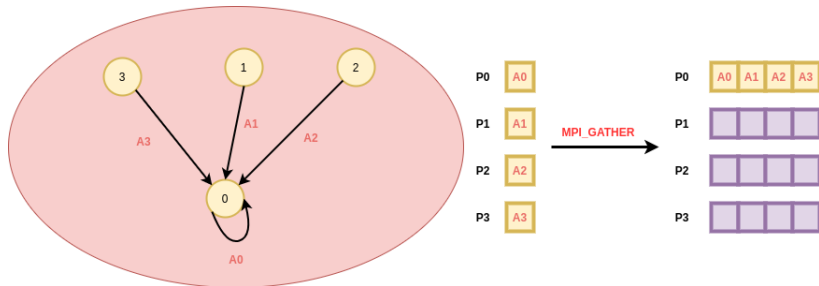


Figure: Collection : MPI_GATHER()

Collective Communications

Global synchronization : MPI_GATHER()

```
1 MPI_GATHER(sendbuf, sendcount, sendtype,  
2           recvbuf, recvcount, recvtype, root, comm, code)  
3  
4 <type> :: sendbuf, recvbuf  
5 integer :: sendcount, recvcount  
6 integer :: sendtype, recvtype  
7 integer :: root, comm, code
```

- Send for each process of communicator comm, a message starting at position sendbuf, of sendcount element type sendtype
- Collect all these messages by the root process at position recvbuf, recvcount element of type recvtype.
- ➡ The couples (sendcount, sendtype) and (recvcount, recvtype) must represent the same size of data.
- ➡ The data are collected in the order of the process ranks.

```
1 Comm.gather(self, sendbuf, root=0)  
2 #or  
3 Comm.Gather(self, sendbuf, recvbuf, int root=0)
```

Collective Communications

Global synchronization : MPI_GATHER() Full example

```
1  from mpi4py import MPI
2  import numpy as np
3
4
5  COMM = MPI.COMM_WORLD
6  SIZE = COMM.Get_size()
7  RANK = COMM.Get_rank()
8
9  sendbuf = (RANK+1)**2
10 print("I am the process {rank}, I send data {data} to 0".format(rank = RANK, ↵
    data=sendbuf))
11
12 recvbuf = COMM.gather(sendbuf, root=0)
13
14 if RANK == 0:
15     for i in range(SIZE):
16         assert recvbuf[i] == (i+1)**2
17
18     print("I am the process 0, I received data {data}".format(data=recvbuf))
```

Collective Communications

Global synchronization : MPI_GATHER() Full example

```
mpirun -n 4 python gather.py
```

```
I am the process 1, I send data 4 to 0
```

```
I am the process 0, I send data 1 to 0
```

```
I am the process 2, I send data 9 to 0
```

```
I am the process 3, I send data 16 to 0
```

```
I am the process 0, I received data [1, 4, 9, 16]
```

Collective Communications

Global synchronization : MPI_ALLGATHER()

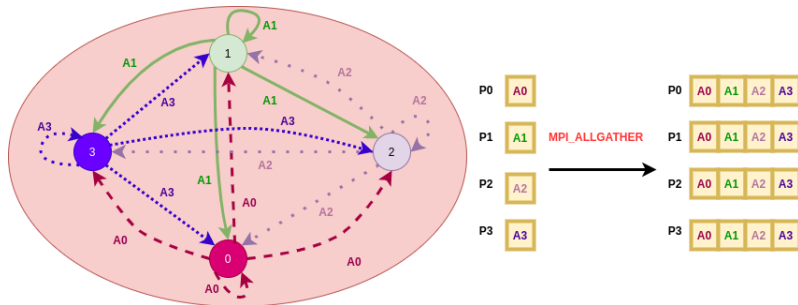


Figure: Gather-to-all : MPI_ALLGATHER()

Collective Communications

Global synchronization : MPI_ALLGATHER()

```
1 MPI_ALLGATHER(sendbuf, sendcount, sendtype,
2               recvbuf, recvcnt, recvtype, comm, code)
3
4 <type> :: sendbuf, recvbuf
5 integer :: sendcount, recvcnt
6 integer :: sendtype, recvtype
7 integer :: comm, code
```

- Corresponds to an MPI_GATHER() followed by an MPI_BCAST()
- Send by each process of communicator comm, a message starting at position sendbuf, of sendcount element, type sendtype.
- Collect all these messages, by all the processes, at position recvbuf of recvcnt element type recvtype.
- ➡ The couples (sendcount, sendtype) and (recvcnt, recvtype) must represent the same data size.
- ➡ The data are gathered in the order of the process ranks.

```
1 Comm.Allgather(self, sendbuf, recvbuf)
2 #or
3 Comm.allgather(self, sendobj)
```

Collective Communications

Global synchronization : MPI_ALLGATHER() Full example

```
1  from mpi4py import MPI
2  import numpy as np
3
4
5  COMM = MPI.COMM_WORLD
6  SIZE = COMM.Get_size()
7  RANK = COMM.Get_rank()
8
9  sendbuf = np.array((RANK+1)**2)
10 print("I am the process {rank}, I send data {data} to all".format(rank = RANK, ↵
    data=sendbuf))
11
12 recvbuf = np.zeros(SIZE, dtype=np.int)
13
14 COMM.Allgather([sendbuf, MPI.INT], [recvbuf, MPI.INT])
15
16
17 print("I am the process {rank}, I received data {data}".format(rank = RANK, ↵
    data=recvbuf))
```

Collective Communications

Global synchronization : `MPI_ALLGATHER()` Full example

```
I am the process 0, I send data 1 to all  
I am the process 3, I send data 16 to all  
I am the process 1, I send data 4 to all  
I am the process 2, I send data 9 to all
```

```
I am the process 1, I received data [ 1  4  9 16]  
I am the process 2, I received data [ 1  4  9 16]  
I am the process 3, I received data [ 1  4  9 16]  
I am the process 0, I received data [ 1  4  9 16]
```

Collective Communications

Global synchronization : `MPI_GATHERV()`

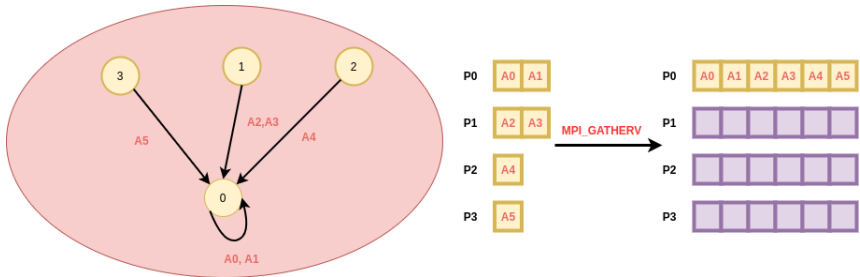


Figure: Extended gather : `MPI_GATHERV()`

Collective Communications

Global synchronization : MPI_GATHERV()

```
1 MPI_GATHERV(sendbuf, sendcount, sendtype,  
2             recvbuf, recvcounts, displs, recvtype,  
3             root, comm, code)  
4  
5 <type> :: sendbuf, recvbuf  
6 integer :: sendcount  
7 integer :: sendtype, recvtype  
8 integer, dimension(:) :: recvcounts, displs  
9 integer :: root, comm, code
```

- This is an MPI_GATHER() where the size of messages can be different among processes
- The i-th process of the communicator comm sends to process root, a message starting at position sendbuf, of sendcount element of type sendtype, and receives at position recvbuf, of recvcounts(i) element of type recvtype, with a displacement of displs(i)
- ➡ The couples (sendcount, sendtype) of the i-th process and (recvcounts(i), recvtype) of process root must be such that the data size sent and received is the same.

```
1 Comm.Gatherv(self, sendbuf, recvbuf, int root=0)
```

Collective Communications

Global synchronization : MPI_GATHERV() Full example

```
1 import numpy as np
2 from mpi4py import MPI
3 import random
4
5 COMM = MPI.COMM_WORLD
6 RANK = COMM.Get_rank()
7
8 local_array = [RANK] * random.randint(2, 5)
9 print("rank: {}, local_array: {}".format(RANK, local_array))
10
11 sendbuf = np.array(local_array)
12
13 # Collect local array sizes using the high-level mpi4py gather
14 sendcounts = np.array(COMM.gather(len(sendbuf), root=0))
15
16 if RANK == 0:
17     print("sendcounts: {}, total: {}".format(sendcounts, sum(sendcounts)))
18     recvbuf = np.empty(sum(sendcounts), dtype=int)
19 else:
20     recvbuf = None
21
22 COMM.Gatherv(sendbuf=sendbuf, recvbuf=(recvbuf, sendcounts), root=0)
23 if RANK == 0:
24     print("Gathered array: {}".format(recvbuf))
```

Collective Communications

Global synchronization : MPI_GATHERV() Full example

```
mpirun -n 4 python gather.py
```

```
rank: 0, local_array: [0, 0, 0, 0, 0]
```

```
rank: 3, local_array: [3, 3, 3, 3, 3]
```

```
rank: 1, local_array: [1, 1, 1]
```

```
rank: 2, local_array: [2, 2, 2]
```

```
sendcounts: [5 3 3 5], total: 16
```

```
Gathered array: [0 0 0 0 0 1 1 1 2 2 2 3 3 3 3 3]
```

Collective Communications

Global synchronization : `MPI_ALLTOALL()`

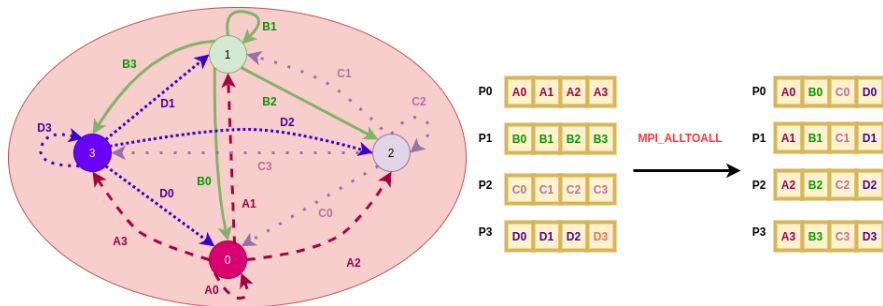


Figure: Collection and distribution : `MPI_ALLTOALL()`

Collective Communications

Global synchronization : MPI_ALLTOALL()

```
1 MPI_ALLTOALL(sendbuf, sendcount, sendtype,
2             recvbuf, recvcount, recvtype, comm, code)
3
4 <type> :: sendbuf, recvbuf
5 integer :: sendcount, recvcount
6 integer :: sendtype, recvtype
7 integer :: comm, code
```

- the i-th process sends its j-th chunk to the j-th process which places it in its i-th chunk.
- ➡ The couples (sendcount, sendtype) and (recvcount, recvtype) must be such that they represent equal data sizes.

```
1 Comm.Alltoall(self, sendbuf, recvbuf)
2 #or
3 Comm.alltoall(self, sendobj)
```

Collective Communications

Global synchronization : MPI_ALLTOALL() Full example

```
1  from mpi4py import MPI
2
3  COMM = MPI.COMM_WORLD
4  RANK = COMM.Get_rank()
5
6  nb_values=4
7  sendbuf = [1000.+RANK*nb_values+i for i in range(nb_values)]
8  print("I am the process {rank}, I send data {data}" .format(rank = RANK, data=↵
      sendbuf))
9
10 recvbuf = COMM.alltoall(sendbuf)
11 print("I am the process {rank}, I received data {data}" .format(rank = RANK, data=↵
      recvbuf))
```

`mpirun -n 4 python alltoall.py`

```
I am the process 1, I send data [1004.0, 1005.0, 1006.0, 1007.0]
I am the process 2, I send data [1008.0, 1009.0, 1010.0, 1011.0]
I am the process 0, I send data [1000.0, 1001.0, 1002.0, 1003.0]
I am the process 3, I send data [1012.0, 1013.0, 1014.0, 1015.0]
```

```
I am the process 2, I received data [1002.0, 1006.0, 1010.0, 1014.0]
I am the process 1, I received data [1001.0, 1005.0, 1009.0, 1013.0]
I am the process 0, I received data [1000.0, 1004.0, 1008.0, 1012.0]
I am the process 3, I received data [1003.0, 1007.0, 1011.0, 1015.0]
```

Collective Communications

Global Reduction

- A reduction is an operation applied to a set of elements in order to obtain one single value. Typical examples are the sum of the elements of a vector ($\text{SUM}(A(:))$) or the search for the maximum value element in a vector ($\text{MAX}(V(:))$).
- MPI proposes high-level subroutines in order to operate reductions on data distributed on a group of processes. The result is obtained on only one process ($\text{MPI_REDUCE}()$) or on all the processes ($\text{MPI_ALLREDUCE}()$), which is in fact equivalent to an $\text{MPI_REDUCE}()$ followed by an $\text{MPI_BCAST}()$).
- If several elements are implied by process, the reduction function is applied to each one of them (for instance to each element of a vector).

Collective Communications

Global Reduction

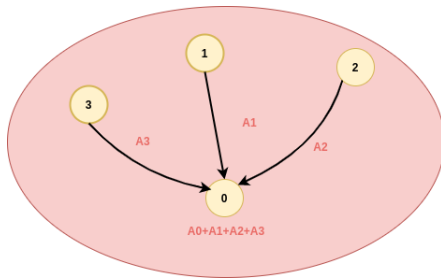


Figure: Distributed reduction (sum)

Collective Communications

Global Reduction: Operations

Name	Operation
MPI_SUM	Sum of elements
MPI_PROD	Product of elements
MPI_MAX	Maximum of elements
MPI_MIN	Minimum of elements
MPI_MAXLOC	Maximum of elements and location
MPI_MINLOC	Minimum of elements and location
MPI_LAND	Logical AND
MPI_LOR	Logical OR
MPI_LXOR	Logical exclusive OR

Table: Global Reduction available operations

Collective Communications

Global Reduction: MPI_REDUCE()

```
1 MPI_REDUCE(sendbuf,recvbuf,count,datatype,op,root,comm,code)
2
3 <type> :: sendbuf, recvbuf
4 integer :: count, datatype, root
5 integer :: op, comm, code
```

- Distributed reduction of count elements of type datatype, starting at position sendbuf, with the operation op from each process of the communicator comm,
- Return the result at position recvbuf in the process root

```
1 Comm.Reduce(self, sendbuf, recvbuf, Op op=SUM, int root=0)
2 #or
3 Comm.reduce(self, sendobj, op=SUM, int root=0)
```

Collective Communications

Global Reduction: MPI_REDUCE() Full example

```
1  from mpi4py import MPI
2  import numpy as np
3
4  COMM = MPI.COMM_WORLD
5  RANK = COMM.Get_rank()
6
7  np.random.seed(RANK)
8
9  data = np.random.randint(2, 10)
10
11  print("The data of rank {rank} is {data}".format(rank=RANK, data=data))
12
13  max_reduced_data = COMM.reduce(data, op=MPI.MAX, root=0)
14  min_reduced_data = COMM.reduce(data, op=MPI.MIN, root=0)
15  sum_reduced_data = COMM.reduce(data, op=MPI.SUM, root=0)
16  prod_reduced_data = COMM.reduce(data, op=MPI.PROD, root=0)
17
18
19  if RANK==0:
20      print("I, process", RANK, "the max of data is :", max_reduced_data)
21      print("I, process", RANK, "the min of data is :", min_reduced_data)
22      print("I, process", RANK, "the sum of data is :", sum_reduced_data)
23      print("I, process", RANK, "the product of data is :", prod_reduced_data)
```

Collective Communications

Global Reduction: MPI_REDUCE() Full example

```
mpirun -n 4 python reduce.py
```

```
The data of rank 0 is 6
```

```
The data of rank 1 is 7
```

```
The data of rank 2 is 2
```

```
The data of rank 3 is 4
```

```
I, process 0 the max of data is : 7
```

```
I, process 0 the min of data is : 2
```

```
I, process 0 the sum of data is : 19
```

```
I, process 0 the product of data is : 336
```


Collective Communications

Global Reduction: `MPI_ALLREDUCE()`

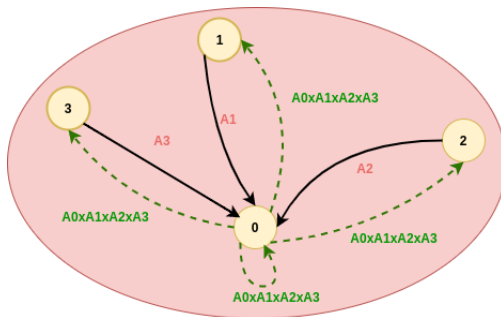


Figure: Distributed reduction (product) with distribution of the result

Collective Communications

Global Reduction: MPI_ALLREDUCE()

```
1 MPI_ALLREDUCE(sendbuf, recvbuf, count, datatype, op, comm, code)
2
3 <type> :: sendbuf, recvbuf
4 integer :: count, datatype
5 integer :: op, comm, code
```

- Distributed reduction of count elements of type datatype starting at position sendbuf, with the operation op from each process of the communicator comm,
- Write the result at position recvbuf for all the processes of the communicator comm.

```
1 Comm.Allreduce(self, sendbuf, recvbuf, Op op=SUM)
2 #or
3 Comm.allreduce(self, sendobj, op=SUM)
```

Collective Communications

Global Reduction: MPI_ALLREDUCE() Full example

```
1  from mpi4py import MPI
2  import numpy as np
3
4  COMM = MPI.COMM_WORLD
5  RANK = COMM.Get_rank()
6
7  np.random.seed(RANK)
8
9  data = np.random.randint(2, 10)
10
11 print("The data of rank {rank} is {data}".format(rank=RANK, data=data))
12
13 max_reduced_data = COMM.allreduce(data, op=MPI.MAX)
14
15 print("I, process {RANK}, The max of data is {max_reduced_data}:".format(RANK=RANK, max_reduced_data=max_reduced_data))
```

Collective Communications

Global Reduction: MPI_ALLREDUCE() Full example

```
mpirun -n 4 python allreduce.py
```

```
The data of rank 1 is 7
```

```
The data of rank 3 is 4
```

```
The data of rank 0 is 6
```

```
The data of rank 2 is 2
```

```
I, process 3, The max of data is 7
```

```
I, process 0, The max of data is 7
```

```
I, process 1, The max of data is 7
```

```
I, process 2, The max of data is 7
```