

# Communicators

## Introduction

The purpose of communicators is to create subgroups on which we can carry out operations such as collective or point-to-point communications. Each subgroup will have its own communication space.

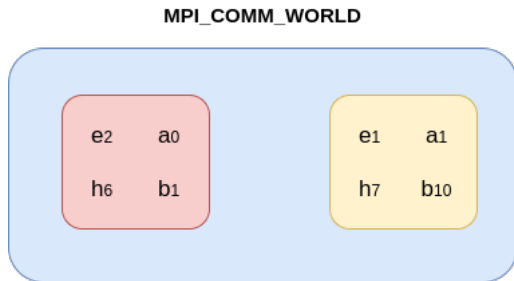


Figure: Communicator partitioning

# Communicators

## Example

For example, we want to broadcast a collective message to even-ranked processes and another message to odd-ranked processes.

- Looping on send/recv can be very detrimental especially if the number of processes is high. Also a test inside the loop would be compulsory in order to know if the sending process must send the message to an even or odd process rank.
- A solution is to create a communicator containing the even-ranked processes, another containing the odd-ranked processes, and initiate the collective communications inside these groups.

# Communicators

## Default communicator

- A communicator can only be created from another communicator. The first one will be created from the `MPI_COMM_WORLD`.
- After the `MPI_INIT()` call, a communicator is created for the duration of the program execution.
- Its identifier `MPI_COMM_WORLD` is an integer value defined in the header files.
- This communicator can only be destroyed via a call to `MPI_FINALIZE()`.
- By default, therefore, it sets the scope of collective and point-to-point communications to include all the processes of the application.

# Communicators

## Groups and communicators

- A communicator consists of :
  - A group, which is an ordered group of processes.
  - A communication context put in place by calling one of the communicator construction subroutines, which allows determination of the communication space.
- The communication contexts are managed by MPI (the programmer has no action on them : It is a hidden attribute).
- In the MPI library, the following subroutines exist for the purpose of building communicators : `MPI_COMM_CREATE()`, `MPI_COMM_DUP()`, `MPI_COMM_SPLIT()`
- The communicator constructors are collective calls.
- Communicators created by the programmer can be destroyed by using the `MPI_COMM_FREE()` subroutine.

# Communicators

## Partitioning of a communicator

In order to solve the problem example :

- Partition the communicator into odd-ranked and even-ranked processes.
- Broadcast a message inside the odd-ranked processes and another message inside the even-ranked processes.




Figure: Communicator creation/destruction

# Communicators

## Partitioning of a communicator with MPI\_COMM\_SPLIT()

The MPI\_COMM\_SPLIT() subroutine allows :

- Partitioning a given communicator into as many communicators as we want.
- Giving the same name to all these communicators : The process value will be the value of its communicator.
- Method :
  1. Define a colour value for each process, associated with its communicator number.
  2. Define a key value for ordering the processes in each communicator
  3. Create the partition where each communicator is called new\_comm

```
1  MPI_COMM_SPLIT(comm,color,key,new_comm,code)
2
3  integer, intent(in) :: comm, color, key
4
5  integer, intent(out) :: new_comm, code
```

```
1  Comm.Split(self, int color=0, int key=0)
```

A process which assigns a color value equal to MPI\_UNDEFINED will have the invalid communicator MPI\_COMM\_NULL for new\_com.

# Communicators

## Example

Let's look at how to proceed in order to build the communicator which will subdivide the communication space into odd-ranked and even-ranked processes via the `MPI_COMM_SPLIT()` constructor.

# Communicators

## Partitioning of a communicator with MPI\_COMM\_SPLIT()

Example :

```
1  ...
2  world_rank = comm.Get_rank(); world_size = comm.Get_size()
3
4  m = 5; a = np.zeros(m)
5
6  if world_rank==2: a[:] = 2.
7  if world_rank==5: a[:] = 5.
8
9  key = world_rank
10 if world_rank==2 or world_rank==5 :
11     key=-1
12 color = world_rank%2
13
14 newcomm = comm.Split(color, key)
15
16 row_rank = newcomm.Get_rank(); row_size = newcomm.Get_size()
17
18 commname = "Comm-"+str(color)
19 newcomm.Set_name(commname)
20
21 print("WORLD RANK/SIZE: {RANK}/{SIZE} \t ROW RANK/SIZE: {ROW_RANK}/{ROW_SIZE} \t ↔
      Comm name: {commname}".format(RANK=world_rank, SIZE=world_size, ROW_RANK=↔
      row_rank, ROW_SIZE=row_size, commname=commname))
22
23 newcomm.Bcast(a, root=0);
24
25 newcomm.Free()
```



# Communicators

## Partitioning of a communicator with MPI\_COMM\_SPLIT()

Results :

```
mpirun -n 6 python3 split_communicator.py
```

WORLD RANK/SIZE: 5/6	ROW RANK/SIZE: 0/3	Comm name: Comm-1
WORLD RANK/SIZE: 0/6	ROW RANK/SIZE: 1/3	Comm name: Comm-0
WORLD RANK/SIZE: 1/6	ROW RANK/SIZE: 1/3	Comm name: Comm-1
WORLD RANK/SIZE: 2/6	ROW RANK/SIZE: 0/3	Comm name: Comm-0
WORLD RANK/SIZE: 3/6	ROW RANK/SIZE: 2/3	Comm name: Comm-1
WORLD RANK/SIZE: 4/6	ROW RANK/SIZE: 2/3	Comm name: Comm-0

# Communicators

## Communicator built from a group

- We can also build a communicator by defining a group of processes : Call to `MPI_COMM_GROUP()`, `MPI_GROUP_INCL()`, `MPI_COMM_CREATE()`, `MPI_GROUP_FREE()`
- This process is however far more cumbersome than using `MPI_COMM_SPLIT()` whenever possible.

# Communicators

## Topologies

- In most applications, especially in domain decomposition methods where we match the calculation domain to the process grid, it is helpful to be able to arrange the processes according to a regular topology.
- MPI allows defining virtual cartesian or graph topologies.
  - Cartesian topologies :
    - | Each process is defined in a grid.
    - | Each process has a neighbour in the grid.
    - | The grid can be periodic or not.
    - | The processes are identified by their coordinates in the grid.
  - Graph topologies :
    - | Can be used in more complex topologies.

# Communicators

## Cartesian topologies

- A Cartesian topology is defined from a given communicator named `comm_old`, calling the `MPI_CART_CREATE()` subroutine.
- We define :
  - An integer `ndims` representing the number of grid dimensions.
  - An integer array `dims` of dimension `ndims` showing the number of processes in each dimension.
  - An array of `ndims` logicals which shows the periodicity of each dimension.
  - A logical `reorder` which shows if the process numbering can be changed by MPI.

```
1 MPI_CART_CREATE(comm_old, ndims, dims, periods, reorder, comm_new, code)
2
3 integer, intent(in) :: comm_old, ndims
4 integer, dimension(ndims), intent(in) :: dims
5 logical, dimension(ndims), intent(in) :: periods
6 logical, intent(in) :: reorganization
7
8 integer, intent(out) :: comm_new, code
```

```
1 Intracomm.Create_cart(self, dims, periods=None, bool reorder=False)
```

# Communicators

## 2D Example

Example on a grid having 2 domains along x and 2 along y, periodic in y.

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 nb_procs = comm.Get_size()
5 rank = comm.Get_rank()
6
7 periods = tuple([False, False])
8 reorder = False
9 dims = [2,2]
10
11
12 cart2d = comm.Create_cart(
13     dims      = dims,
14     periods   = periods,
15     reorder   = reorder
16 )
```

- If `reorder = .false.` then the rank of the processes in the new communicator (`comm_2D`) is the same as in the old communicator (`MPI.COMM_WORLD`).
- If `reorder = .true.`, the MPI implementation chooses the order of the processes.




Figure: A 2D non-periodic Cartesian topology

# Communicators

## 3D Example

Example on a 3D grid having 2 domains along x, 2 along y and 2 along z, non periodic.

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 nb_procs = comm.Get_size()
5 rank = comm.Get_rank()
6
7 periods = tuple([False, False, False])
8 reorder = False
9 dims = [2,2,2]
10
11
12 cart3 = comm.Create_cart(
13     dims = dims,
14     periods = periods,
15     reorder = reorder
16 )
```





Figure: A 3D non-periodic Cartesian topology



# Communicators

## Process distribution

The `MPI_DIMS_CREATE()` subroutine returns the number of processes in each dimension of the grid according to the total number of processes.

```
1 MPI_DIMS_CREATE(nb_procs, ndims, dims, code)
2
3 integer, intent(in) :: nb_procs, ndims
4 integer, dimension(ndims), intent(inout) :: dims
5
6 integer, intent(out) :: code
```

Remark : If the values of `dims` in entry are all 0, then we leave to MPI the choice of the number of processes in each direction according to the total number of processes.

# Communicators

## Rank of a process

In a Cartesian topology, the `MPI_CART_RANK()` subroutine returns the rank of the associated process to the coordinates in the grid.

```
1 MPI_CART_RANK(comm, coords, rank, code)
2
3 integer, intent(in) :: comm
4 integer, dimension(ndims), intent(in) :: coords
5
6 integer, intent(out) :: rank, code
```

```
1 Cartcomm.Get_cart_rank(self, coords)
```

# Communicators

## Coordinates of a process

In a cartesian topology, the `MPI_CART_COORDS()` subroutine returns the coordinates of a process of a given rank in the grid.

```
1 MPI_CART_COORDS(comm, rank, ndims, coords, code)
2
3 integer, intent(in) :: comm, rank, ndims
4 integer, dimension(ndims), intent(out) :: coords
5
6 integer, intent(out) :: code
```

```
1 Cartcomm.Get_coords(self, int rank)
```

# Communicators

Example : `MPI_CART_COORDS()`




Figure: A 2D non-periodic Cartesian topology

# Communicators

## Example : MPI\_CART\_COORDS()

```
1  from mpi4py import MPI
2
3  comm = MPI.COMM_WORLD
4  rank = comm.Get_rank()
5
6  periods = tuple([False, False])
7  reorder = False
8  dims = [2, 2]
9
10 cart2d = comm.Create_cart(
11     dims      = dims,
12     periods   = periods,
13     reorder   = reorder
14 )
15
16 coord2d = cart2d.Get_coords(rank)
17
18 print("I'm rank", rank, "my 2d coords are", coord2d)
```

`mpirun -n 4 python3 coordinate_2d_cart.py`

I'm rank 0 my 2d coords are [0, 0]

I'm rank 1 my 2d coords are [0, 1]

I'm rank 2 my 2d coords are [1, 0]

I'm rank 3 my 2d coords are [1, 1]

# Communicators

## Rank of neighbours

In a Cartesian topology, a process that calls the `MPI_CART_SHIFT()` subroutine can obtain the rank of a neighboring process in a given direction.

```
1 MPI_CART_SHIFT(comm, direction, step, rank_previous, rank_next, code)
2
3 integer, intent(in) :: comm, direction, step
4 integer, intent(out) :: rank_previous, rank_next
5
6 integer, intent(out) :: code
```

```
1 Cartcomm.Shift(self, int direction, int disp)
```

- The direction parameter corresponds to the displacement axis (xyz).
- The step parameter corresponds to the displacement step.
- If a rank does not have a neighbor before (or after) in the requested direction, then the value of the previous (or following) rank will be `MPI_PROC_NULL`.

# Communicators

Example : MPI\_CART\_SHIFT()

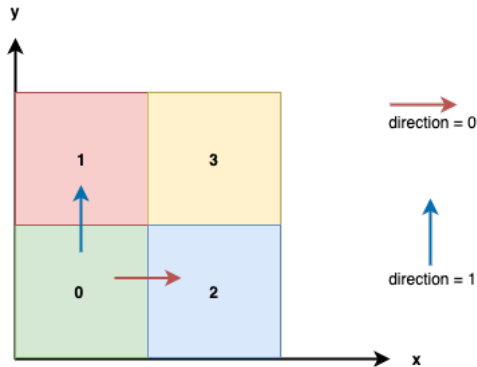


Figure: Call of the MPI\_CART\_SHIFT() subroutine

# Communicators

Example : MPI\_CART\_SHIFT()

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 rank = comm.Get_rank()
5
6 periods = tuple([False, False])
7 reorder = False
8 dims = [2, 2]
9
10 cart2d = comm.Create_cart(
11     dims = dims,
12     periods = periods,
13     reorder = reorder
14 )
15
16 left, right = cart2d.Shift(direction = 0, disp=1)
17 low, high = cart2d.Shift(direction = 1, disp=1)
18
19 print("I'm rank", rank, "my (left, right) neighbours are", (left, right),
20       "my (low, high) neighbours are", (low, high))
```



# Communicators

Example : MPI\_CART\_SHIFT()

```
mpirun -n 4 python3 neighbours_2d_cart.py
```

```
I'm rank 0
```

```
my (left,right) neighbours are (-2, 2) my (low,high) neighbours are (-2, 1)
```

```
I'm rank 1
```

```
my (left,right) neighbours are (-2, 3) my (low,high) neighbours are (0, -2)
```

```
I'm rank 2
```

```
my (left,right) neighbours are (0, -2) my (low,high) neighbours are (-2, 3)
```

```
I'm rank 3
```

```
my (left,right) neighbours are (1, -2) my (low,high) neighbours are (2, -2)
```

# Communicators

## 3D Example : coordinates and neighbours




Figure: 3D Coordinates and Neighbours

# Communicators

## 3D Example : coordinates and neighbours

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 rank = comm.Get_rank()
5
6 periods = tuple([False, False, False])
7 reorder = False
8 dims = [2,2,2]
9
10 cart3d = comm.Create_cart(
11     dims      = dims,
12     periods   = periods,
13     reorder   = reorder
14 )
15
16 coord3d      = cart3d.Get_coords(rank)
17 left,right   = cart3d.Shift(direction = 0, disp=1)
18 low,high     = cart3d.Shift(direction = 1, disp=1)
19 ahead,before = cart3d.Shift(direction = 2, disp=1)
20
21 print("I'm rank", rank, "my 3d coords are", coord3d, "my (left,right) neighbours ↔
      are", (left, right), "my (low,high) neighbours are", (low, high), "my (ahead↔
      ,before) neighbours are", (ahead,before))
```

# Communicators

## 3D Example : coordinates and neighbours

```
mpirun -n 8 --oversubscribe python3 create_3d_cart.py
```

```
I'm rank 0 my 3d coords are [0, 0, 0]
my (left,right) neighbours are (-2, 4)
my (low,high) neighbours are (-2, 2)
my (ahead,before) neighbours are (-2, 1)
```

```
I'm rank 1 my 3d coords are [0, 0, 1]
my (left,right) neighbours are (-2, 5)
my (low,high) neighbours are (-2, 3)
my (ahead,before) neighbours are (0, -2)
```

```
I'm rank 2 my 3d coords are [0, 1, 0]
my (left,right) neighbours are (-2, 6)
my (low,high) neighbours are (0, -2)
my (ahead,before) neighbours are (-2, 3)
```

```
...
```